

# Introspection via Self Debugging

Russell Harmon  
`reh5586@cs.rit.edu`

Rochester Institute of Technology  
Computer Science

December 12, 2013

# Overview

Debuggers

Introspection

Introspection + Debugging

Ruminate

Use

Internals

Limitations

Future Work

# Debugger

*... a computer program that is used to test and debug other programs -Wikipedia*

# Debugger

*... a computer program that is used to test and debug other programs -Wikipedia*

- ▶ LLDB - LLVM debugger
- ▶ GDB - GNU debugger
- ▶ WinDBG - Windows debugger
- ▶ DBX - Sun Studio debugger
- ▶ LDB - Intel debugger

## Debugging Example

```
1  struct foo {  
2      char *str;  
3  };  
4  int main() {  
5      struct foo bar;  
6      bar.str = "Hello World!";  
7      asm("int $0x3");  
8  }
```

## Debugging Example

```
(lldb) run
* thread #1: tid = 0x1c03, 0x00000000100000f65
  a.out`main + 21 at struct.c:9, stop reason =
  EXC_BREAKPOINT (code=EXC_I386_BPT, subcode=0x0)
    frame #0: 0x00000000100000f65 a.out`main + 21
    at struct.c:9
      6      struct foo bar;
      7      bar.str = "Hello World!";
      8      asm("int $0x3");
-> 9      }
```

## Debugging Example

```
(lldb) run
* thread #1: tid = 0x1c03, 0x00000000100000f65
  a.out`main + 21 at struct.c:9, stop reason =
  EXC_BREAKPOINT (code=EXC_I386_BPT, subcode=0x0)
    frame #0: 0x00000000100000f65 a.out`main + 21
    at struct.c:9
      6      struct foo bar;
      7      bar.str = "Hello World!";
      8      asm("int $0x3");
-> 9      }
(lldb) print bar
(foo) $0 = {
  (char *) str = 0x00000000100000f67 "Hello World!"
}
```

## Debugging Example

```
(lldb) print bar
(foo) $0 = {
    (char *) str = 0x00000000100000f67 "Hello World!"
}
```



## Debugging Example

```
(lldb) print bar
(foo) $0 = {
    (char *) str = 0x00000001000000f67 "Hello World!"
}
```

```
1 struct foo {
2     char *str;
3 };
4 int main() {
5     struct foo bar;
6     bar.str = "Hello World!";
7     asm("int $0x3");
8 }
```

# Introspection

*... the ability for a program to examine the type or properties of an object at runtime. -Wikipedia*

- ▶ Interpreted languages have reflection / introspection.
- ▶ introspection  $\subset$  reflection
- ▶ Java: `this.getClass().getCanonicalName()`
- ▶ Ruby: `self.class.name`

## Introspection + Debugging

```
1 struct foo {  
2     char *str;  
3 };  
4 int main() {  
5     struct foo bar;  
6     bar.str = "Hello World!";  
7     printf("%s\n", gettype(bar)->name);  
8     printf("%s\n", gettype(bar.str)->name);  
9 }
```

# Introspection + Debugging

```
1 struct foo {  
2     char *str;  
3 };  
4 int main() {  
5     struct foo bar;  
6     bar.str = "Hello World!";  
7     printf("%s\n", gettype(bar)->name);  
8     printf("%s\n", gettype(bar.str)->name);  
9 }
```

```
foo  
char *
```

# Use Cases

- ▶ `abort_with_stacktrace();`
- ▶ easy data structure traversal
- ▶ access to third party library internals

# Use Cases

- ▶ `abort_with_stacktrace();`
- ▶ easy data structure traversal
- ▶ access to third party library internals

# Use Cases

- ▶ `abort_with_stacktrace();`
- ▶ easy data structure traversal
- ▶ access to third party library internals

# Ruminate

- ▶ Introspective library for C
- ▶ Data introspection (structs, unions, enums, functions, primitives)
- ▶ Stack introspection



# Ruminate

- ▶ Introspective library for C
- ▶ Data introspection (structs, unions, enums, functions, primitives)
- ▶ Stack introspection

# Ruminate

- ▶ Introspective library for C
- ▶ Data introspection (structs, unions, enums, functions, primitives)
- ▶ Stack introspection

# Ruminate

## Simple Use

```
1  // Includes go here
2  struct foo {
3      char *str;
4  };
5  int main( int argc, char *argv[] ) {
6      struct foo bar;
7      bar.str = "Hello World!";
8      ruminate_init(argv[0], NULL);
9      printf("%s\n", r_string_bytes(r_type_name(
10         ruminate_get_type(bar, NULL), NULL)));
11     printf("%s\n", r_string_bytes(r_type_name(
12         ruminate_get_type(bar.str, NULL), NULL)));
13 }
```

# Ruminate

## Simple Use

```
1  // Includes go here
2  struct foo {
3      char *str;
4  };
5  int main( int argc, char *argv[] ) {
6      struct foo bar;
7      bar.str = "Hello World!";
8      ruminate_init(argv[0], NULL);
9      printf("%s\n", r_string_bytes(r_type_name(
10         ruminate_get_type(bar, NULL), NULL)));
11     printf("%s\n", r_string_bytes(r_type_name(
12         ruminate_get_type(bar.str, NULL), NULL)));
13 }

foo
char *
```

# Ruminate

## Simple Use

A rundown of the ruminare functions used

- ▶ `ruminate_init` initializes `Ruminate`
- ▶ `ruminate_get_type` retrieves an `RType` representing a type
- ▶ `r_type_name` retrieves an `RString` containing the name of an `RType`
- ▶ `r_string_bytes` retrieves a `char *` containing the value of an `RString`

# Ruminate

Ok, but that's really simple...

# Ruminate

## Type Traversal

```
1  // Includes go here
2  struct foo { char *str; };
3  int main( int argc, char *argv[] ) {
4      struct foo bar;
5      bar.str = "Hello World!";
6      ruminate_init(argv[0], NULL);
7      RType *type = ruminate_get_type(bar, NULL);
8      RAggregateType *agg_type = (RAggregateType *) type;
9      RAggregateMember *type_member =
10         r_aggregate_type_member_by_name(agg_type, "str", NULL);
11      RType *member_type = r_type_member_type(
12         (RTypeMember *) type_member, NULL);
13      RString *member_name = r_type_name(member_type, NULL);
14      const char *member_name_str = r_string_bytes(member_name);
15      printf("%s\n", member_name_str);
16  }
```

# Ruminate

## Type Traversal

```
1  // Includes go here
2  struct foo { char *str; };
3  int main( int argc, char *argv[] ) {
4      struct foo bar;
5      bar.str = "Hello World!";
6      ruminate_init(argv[0], NULL);
7      RType *type = ruminate_get_type(bar, NULL);
8      RAggregateType *agg_type = (RAggregateType *) type;
9      RAggregateMember *type_member =
10         r_aggregate_type_member_by_name(agg_type, "str", NULL);
11      RType *member_type = r_type_member_type(
12         (RTypeMember *) type_member, NULL);
13      RString *member_name = r_type_name(member_type, NULL);
14      const char *member_name_str = r_string_bytes(member_name);
15      printf("%s\n", member_name_str);
16  }
```

char \*



# Ruminate


## Type Traversal

Only a few new functions

- ▶ `r_aggregate_type_member_by_name` gets a member of an aggregate type by it's name
- ▶ `r_type_member_type` gets the type of an aggregate member

# Ruminate

## Type Hierarchy



# Ruminate

Still too simple?

# Ruminate

Still too simple?

Ok, well there's this

# Ruminate

## Complex Example

```
#include <ruminate.h>
#include <stdlib.h>
#include <string.h>
#include <stdio.h>

typedef char *string_t;

struct foo {
    string_t str;
    int i;
};

void print_data( RType *type, const void *data ) {
    switch( r_type_id(type, NULL) ) {
        case R_TYPE_TYPEDEF:
            if( strcmp(r_string_bytes(r_type_name(type, NULL)), "string_t") == 0 )
                printf("(string_t) \"%s\"\n", *((const string_t *) data));
            break;
        case R_TYPE_BUILTIN:
            switch( r_builtin_type_id((RBuiltinType *) type, NULL) ) {
                case R_BUILTIN_TYPE_INT:
                    printf("(int) %d\n", *((const int *) data));
                    break;
            }
            break;
        case R_TYPEAggregate: {
            RAggregateType *agg = (RAggregateType *) type;
            if( r_aggregate_type_id(agg, NULL) == RAggregateType_STRUCT ) {
                printf("(%)s {\n", r_string_bytes(r_type_name(type, NULL)));
                for( size_t i = 0; i < r_aggregate_type_nmembers(agg, NULL); i++ ) {
                    RAggregateMember *memb = r_aggregate_type_member_at(agg, i, NULL);
                    printf("\t.%s = ", r_string_bytes(r_aggregate_member_name(memb, NULL)));
                    RTypeMember *tmemb = (RTypeMember *) memb;
                    off_t offset = r_type_member_offset(tmemb, NULL);
                    print_data(r_type_member_type(tmemb, NULL), data + offset);
                }
                printf("}\n");
            }
            break;
        }
    }
}

int main( int argc, char *argv[] ) {
    (void) argc;
    ruminate_init(argv[0], NULL);
    struct foo bar = {
        .str = "Hello World!",
        .i = 6666
    };
    print_data(ruminate_get_type(bar, NULL), &bar);
}
```

# Ruminate

But that's too large to fit.

Ruminate

Oh wait

# Ruminate

Oh wait

Lets not forget



# Ruminate

Oh wait

Lets not forget

You can do stack introspection too

# Ruminate

## Stack Introspection

```
1 // Includes go here
2 void print_backtrace() {
3     RFrameList *frames = ruminde_backtrace(NULL);
4     size_t frames_len = r_frame_list_size(frames, NULL);
5     for( size_t i = 0; i < frames_len; i++ )
6         printf("%s\n", r_string_bytes(r_frame_function_name(
7             r_frame_list_at(frames, i, NULL), NULL)));
8 }
9 void bar( int i ) { if( i < 2 ) bar(i + 1); else print_backtrace(); }
10 void foo() { bar(0); }
11 int main( int argc, char *argv[] ) {
12     ruminde_init(argv[0], NULL);
13     foo();
14 }
```

# Ruminate

## Stack Introspection

```
1 // Includes go here
2 void print_backtrace() {
3     RFrameList *frames = ruminde_backtrace(NULL);
4     size_t frames_len = r_frame_list_size(frames, NULL);
5     for( size_t i = 0; i < frames_len; i++ )
6         printf("%s\n", r_string_bytes(r_frame_function_name(
7             r_frame_list_at(frames, i, NULL), NULL)));
8 }
9 void bar( int i ) { if( i < 2 ) bar(i + 1); else print_backtrace(); }
10 void foo() { bar(0); }
11 int main( int argc, char *argv[] ) {
12     ruminde_init(argv[0], NULL);
13     foo();
14 }
```

```
ruminde_hit_breakpoint
ruminde_backtrace
print_backtrace
bar
bar
bar
foo
main
__libc_start_main
```

# Ruminate

## Typed Memory Allocator

Building on this, you can do typed dynamic memory allocation.

# Ruminate

## Typed Memory Allocator

Building on this, you can do typed dynamic memory allocation.

```
1  // Includes go here
2  int main( int argc, char *argv[] ) {
3      ruminate_init(argv[0], NULL);
4      const char *src_str = "Hello World!";
5      size_t src_str_len = strlen(src_str) + 1;
6      void *str = r_mem_malloc_sized(
7          char *, src_str_len, NULL
8      );
9      memcpy(str, src_str, src_str_len);
10     printf("(s) \"%s\"\\n", r_string_bytes(r_type_name(
11         r_mem_type(str),
12         NULL
13     )), str);
14 }
```

(char \*) "Hello World!"

# Ruminate

## JSON

and do transparent json conversion

# Ruminate

## JSON

and do transparent json conversion  
both unidirectionally

```
1 // Includes go here
2 struct MyStruct {
3     int i;
4     enum MyEnum { MY_ENUM_VALUE_1, MY_ENUM_VALUE_2 } e;
5 };
6 int main( int argc, char *argv[] ) {
7     ruminate_init(argv[0], NULL);
8     struct MyStruct foo = { 1, MY_ENUM_VALUE_2 };
9     json_dumpf(
10         json_serialize(NULL, ruminate_get_type(foo, NULL),
11                        &foo, NULL),
12         stdout, 0
13     );
14 }
```

{"i": 1, "e": 1}

# Ruminate

## JSON

and bidirectionally

```
1 // Includes go here
2 struct MyStruct {
3     int i;
4     enum MyEnum { MY_ENUM_VALUE_1, MY_ENUM_VALUE_2 } e;
5 };
6 int main( int argc, char *argv[] ) {
7     ruminate_init(argv[0], NULL);
8     struct MyStruct foo = { 1, MY_ENUM_VALUE_2 };
9     JsonState *js = json_state_new();
10    json_state_set_flags(js, JSON_FLAG_INVERTABLE);
11    json_dumpf(
12        json_serialize(js, ruminate_get_type(foo, NULL),
13                        &foo, NULL),
14        stdout, 0
15    );
16 }
```

{"value": {"i": 1, "e": 1}, "type": "MyStruct"}



# Implementation

Lets dive in

# Implementation

## RType

```
1  struct RType {
2      RTypeId id;
3      gint refcnt;
4      RString *name;
5      RuminantBackend::TypePrx type;
6      RuminantBackend::TypeId type_id;
7      struct {
8          size_t value;
9          bool initialized;
10     } size;
11 };
```

# Implementation

## RType

```
1  struct RType {
2      RTypeId id;
3      gint refcnt;
4      RString *name;
5      RuminantBackend::TypePrx type;
6      RuminantBackend::TypeId type_id;
7      struct {
8          size_t value;
9          bool initialized;
10     } size;
11 };
```

Aside from caching, RType is just a proxy backed by  
RuminantBackend::TypePrx

# RPC Contract

- ▶ Simple, minimal RPC interface
- ▶ Built off the Type interface
- ▶ Should be simple enough to port to other debuggers

# RPC Contract

- ▶ Simple, minimal RPC interface
- ▶ Built off the Type interface
- ▶ Should be simple enough to port to other debuggers

```
1  interface Type {
2      TypeId getId();
3      Type *getPointeeType();
4      Type *getPointerType();
5      Type *getCanonicalType();
6      Type *getReturnType();
7      Type *getArrayMemberType();
8      idempotent long getArrayLength();
9      TypeMemberList getMembers( optional(1) long tid );
10     string getName();
11     idempotent long getSize();
12     idempotent bool isSigned();
13     idempotent bool isUnsigned();
14 }
```

# Implementation

## RType

RType and it's decendants:

- ▶ type safe
- ▶ caching
- ▶ reference counted
- ▶ wrapper around Type

# Implementation

## Debugger Server

Manipulate LLDB to implement the RPC contract

# Implementation

## Debugger Server

Manipulate LLDB to implement the RPC contract

Some interesting bits

- ▶ Enums
- ▶ Signals
- ▶ Arrays



# Implementation

## LLDB

- ▶ DWARF
- ▶ ELF, Mach-O
- ▶ Name demangling
- ▶ Stack traversal

# Implementation

## Enums

LLDB's public API did not support enum introspection. Support was added.

# Implementation

## Enums

LLDB's public API did not support enum introspection. Support was added.

Enum members are interesting

# Implementation

## Enums

LLDB's public API did not support enum introspection. Support was added.

Enum members are interesting  
the only type in C whose value is part of the type

```
enum MyEnum {  
    MY_ENUM_1,  
    MY_ENUM_2  
};
```

# Implementation

## Signals

*While being traced, the tracee will stop each time a signal is delivered -ptrace(2)*

LLDB default behavior is to stop on (most) signals and wait for user input...

# Implementation

## Signals

*While being traced, the tracee will stop each time a signal is delivered -ptrace(2)*

LLDB default behavior is to stop on (most) signals and wait for user input...

there is no user

# Implementation

## Signal Deadlock

If a signal occurs, LLDB will wait for the debuggee to provide instructions. The debuggee will remain stopped.

# Implementation

## Signal Deadlock

If a signal occurs, LLDB will wait for the debuggee to provide instructions. The debuggee will remain stopped.

change LLDB's signal disposition to ignore signals

No support for this in LLDB. Support was added.



# Implementation

## Arrays

LLDB does not model arrays in it's pure type representation.

# Implementation

## Arrays

LLDB does not model arrays in it's pure type representation.  
Arrays are accessed by value only.

# Implementation

## Arrays

LLDB does not model arrays in it's pure type representation.  
Arrays are accessed by value only.

```
char a[sizeof((*((int (*)[4]) NULL))[0])];  
__typeof__(&a) ap = &a;  
(int (*)[4]) ap
```

# Implementation

## Typed Memory

Typed memory is implemented by padding the beginning of allocated memory with a pointer to the type.

# Implementation

## JSON

Conversion to JSON is simple

```
if the type is a struct
    generate a JSON object
    for each subtype
        recurse
```

```
else if the type is an enum
    generate a JSON integer
```

```
else if the type is an array
    generate a JSON array
    recurse
```

```
else if the type is a pointer
    recurse using the dereferenced pointer
```

```
else
    generate a JSON primitive
```

Note that unions and strings are not handled.

# Implementation

## JSON

Conversion from json is similar

retrieve the name of the type from the top-level json object

lookup the type in the current program

allocate typed memory

proc repeat

  if the type is a struct

    retrieve the JSON object representing this value

    for each subtype

      call repeat

  else if the type is an enum

    insert the value of the JSON int

  else if the type is an array

    retrieve the JSON array representing this value

    call repeat

  else if the type is a pointer

    allocate typed memory

    call repeat

  else

    insert the value from the JSON primitive

# Limitations

- ▶ strings are not detectable
- ▶ the current type of a union is not detectable
- ▶ slow
- ▶ debugging symbols are required

# Limitations

## Strings

strings are not detectable



# Limitations

## Strings

strings are not detectable

In C, a string has type `char *`. That is, it's type is "pointer to byte".

With no additional information it is unknowable if this `char *` type is

- ▶ a valid NULL terminated string
- ▶ a pointer to a single character
- ▶ a pointer to multiple characters which is non NULL terminated
- ▶ a pointer to one or more bytes which are not valid characters in the current encoding

# Limitations

## Unions

- The current type of a union is not detectable

# Limitations

## Unions

The current type of a union is not detectable

```
1  union foo { int a; char b; }
2  int main( int argc ) {
3      union foo f;
4      if( argc > 2 ) {
5          f.a = 1;
6      } else {
7          f.b = 'c';
8      }
9      // !!
10 }
```

- ▶ The compiler cannot know the type of `f` at line 9 due to the indeterminate nature of `argc`
- ▶ The debugger cannot know the type of `f` at line 9 because DWARF does not (and cannot) contain enough information for the debugger to determine the current type of `f`

# Limitations

## Slow

Ruminate imposes significant overhead.

```
// Includes go here
```

```
int main( int argc, char *argv[] ) {  
    ruminate_init(argv[0], NULL);  
    const char *src_str = "Hello World!";  
    size_t src_str_len = strlen(src_str) + 1;  
    void *str = r_mem_malloc_sized(  
        char *, src_str_len, NULL  
    );  
    memcpy(str, src_str, src_str_len);  
    printf("(%)s \"s\"\\n", r_string_bytes(r_type_name(  
        r_mem_type(str),  
        NULL  
    )), str);  
}
```

Takes 1.6 seconds to run.

# Limitations

## Debugging Symbols

Debugging symbols are required

# Limitations

## Debugging Symbols

Debugging symbols are required

Not a limitation

# Limitations

## Debugging Symbols

Debugging symbols are required

Not a limitation

DWARF debugging symbols are made up of two components:

- ▶ line number information
- ▶ type information

all of which is available in every other language that has introspection.

# Limitations

## Debugging Symbols

Debugging symbols are required

Not a limitation

DWARF debugging symbols are made up of two components:

- ▶ line number information
- ▶ type information

all of which is available in every other language that has introspection.

Debugging symbols are only required in modules being introspected.



# Future Work

slow is a problem

## Future Work

slow is a problem

eliminate the debugger controller

# Future Work

slow is a problem

eliminate the debugger controller

embed the debugger inside the debuggee

# Future Work

slow is a problem

eliminate the debugger controller

embed the debugger inside the debuggee

how?

## Future Work

what makes a debugger?

# Future Work

what makes a debugger?

- ▶ remote process controller
- ▶ remote memory inspector
- ▶ executable parser (ELF, Mach-O, etc)
- ▶ symbol parser
- ▶ stack traverser

# Future Work

what does Ruminant need?

# Future Work

what does Ruminator need?

- ▶ remote process controller
- ▶ remote memory inspector
- ▶ executable parser (ELF, Mach-O, etc)
- ▶ symbol parser
- ▶ stack traverser



# Future Work

how?

# Future Work

how?

split LLDB's core into several components

- ▶ symbol parser
- ▶ stack traverser

useful for more than just Ruminator

Thanks